

In [1]:

```

# Import Libraries
import numpy as np # for linear algebra
import pandas as pd # for data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns # for data visualization
import matplotlib.pyplot as plt # to plot data visualization charts
from collections import Counter
import os

# Modeling Libraries
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score
from sklearn.preprocessing import QuantileTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score, StratifiedKFold, learning_curve
from sklearn.svm import SVC

```

In [4]:

```

# Importing the dataset from Kaggle
data = pd.read_csv("diabetes.csv")

# First step is getting familiar with the structure of the dataset
data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Pregnancies           768 non-null   int64  
 1   Glucose               768 non-null   int64  
 2   BloodPressure         768 non-null   int64  
 3   SkinThickness         768 non-null   int64  
 4   Insulin               768 non-null   int64  
 5   BMI                   768 non-null   float64 
 6   DiabetesPedigreeFunction 768 non-null   float64 
 7   Age                  768 non-null   int64  
 8   Outcome               768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

In [5]:

```
data.head()
```

Out[5]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148	72	35	0	33.6	0
1	1	85	66	29	0	26.6	0
2	8	183	64	0	0	23.3	0
3	1	89	66	23	94	28.1	0
4	0	137	40	35	168	43.1	2

Filling the Missing Values

In [6]:

```
data.isnull().sum()
```

Out[6]:

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

In [7]:

```

data['Glucose'] = data['Glucose'].replace(0, data['Glucose'].median())

# Filling 0 values of Blood Pressure
data['BloodPressure'] = data['BloodPressure'].replace(0, data['BloodPressure'].median())

# Replacing 0 values in BMI
data['BMI'] = data['BMI'].replace(0, data['BMI'].mean())

# Replacing the missing values of Insulin and SkinThickness
data['SkinThickness'] = data['SkinThickness'].replace(0, data['SkinThickness'].mean())
data['Insulin'] = data['Insulin'].replace(0, data['Insulin'].mean())
data.head()

```

Out[7]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35.000000	79.799479	33.6	
1	1	85	66	29.000000	79.799479	26.6	
2	8	183	64	20.536458	79.799479	23.3	
3	1	89	66	23.000000	94.000000	28.1	
4	0	137	40	35.000000	168.000000	43.1	

In [8]:

```
data.describe()
```

Out[8]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Dia
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	121.656250	72.386719	26.606479	118.660163	32.450805	
std	3.369578	30.438286	12.096642	9.631241	93.080358	6.875374	
min	0.000000	44.000000	24.000000	7.000000	14.000000	18.200000	
25%	1.000000	99.750000	64.000000	20.536458	79.799479	27.500000	
50%	3.000000	117.000000	72.000000	23.000000	79.799479	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

In []:


```
#Exploratory Data Analysis
```

In [9]:

```
plt.figure(figsize = (10, 8))
sns.heatmap(data.corr(), annot = True, fmt = ".3f", cmap = "YlGnBu")
plt.title("Correlation heatmap")
```

Out[9]:

Text(0.5, 1.0, 'Correlation heatmap')



In [10]:

```
plt.figure(figsize = (10, 8))

# Plotting density function graph of the pregnancies and the target variable
kde = sns.kdeplot(data["Pregnancies"][data["Outcome"] == 1], color = "Red", shade = True)
kde = sns.kdeplot(data["Pregnancies"][data["Outcome"] == 0], ax = kde, color = "Blue", shade = True)
kde.set_xlabel("Pregnancies")
kde.set_ylabel("Density")
kde.legend(["Positive Result", "Negative Result"])
```

C:\Users\venut\AppData\Local\Temp\ipykernel_16960\827401410.py:4: FutureWarning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`. This will become an error in seaborn v0.14.0; please update your code.


```
kde = sns.kdeplot(data["Pregnancies"][data["Outcome"] == 1], color = "Red", shade = True)
C:\Users\venut\AppData\Local\Temp\ipykernel_16960\827401410.py:5: FutureWarning:
```

`shade` is now deprecated in favor of `fill`; setting `fill=True`. This will become an error in seaborn v0.14.0; please update your code.

```
kde = sns.kdeplot(data["Pregnancies"][data["Outcome"] == 0], ax = kde, color = "Blue", shade = True)
```

Out[10]:

<matplotlib.legend.Legend at 0x273f0fbbd60>




In [11]:

```
# Exploring the Glucose and the Target variables together
plt.figure(figsize = (10, 8))
sns.violinplot(data = data, x = "Outcome", y = "Glucose",
               split = True, inner = "quart", linewidth = 2)
```

Out[11]:

<Axes: xlabel='Outcome', ylabel='Glucose'>



In [12]:

```
plt.figure(figsize = (10, 8))
kde = sns.kdeplot(data["Glucose"][data["Outcome"] == 1], color = "Red", shade = True)
kde = sns.kdeplot(data["Glucose"][data["Outcome"] == 0], ax = kde, color = "Blue", shade = True)
kde.set_xlabel("Glucose")
kde.set_ylabel("Density")
kde.legend(["Positive Result", "Negative Result"])
```

C:\Users\venut\AppData\Local\Temp\ipykernel_16960\2458365466.py:2: Future Warning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
kde = sns.kdeplot(data["Glucose"][data["Outcome"] == 1], color = "Red",  
shade = True)
```


C:\Users\venut\AppData\Local\Temp\ipykernel_16960\2458365466.py:3: Future Warning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
kde = sns.kdeplot(data["Glucose"][data["Outcome"] == 0], ax = kde, color = "Blue",  
fill = True)
```

Out[12]:

<matplotlib.legend.Legend at 0x273f0600be0>



In []:

#Implementing Machine Learning Models

In [13]:

```
# Transforming the data into quartiles
quartile = QuantileTransformer()
X = quartile.fit_transform(data)
dataset = quartile.transform(X)
dataset = pd.DataFrame(X)
dataset.columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
# Showing the top 5 rows of the transformed dataset
dataset.head()
```

C:\Users\venut\anaconda3\lib\site-packages\sklearn\preprocessing_data.p
y:2627: UserWarning: n_quantiles (1000) is greater than the total number
of samples (768). n_quantiles is set to n_samples.
warnings.warn(
C:\Users\venut\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarni
ng: X does not have valid feature names, but QuantileTransformer was fitt
ed with feature names
warnings.warn(

Out[13]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigri
0	0.747718	0.810300	0.494133	0.801825	0.380052	0.591265	
1	0.232725	0.091265	0.290091	0.644720	0.380052	0.213168	
2	0.863755	0.956975	0.233377	0.308996	0.380052	0.077575	
3	0.232725	0.124511	0.290091	0.505867	0.662973	0.284224	
4	0.000000	0.721643	0.005215	0.801825	0.834420	0.926988	

In []:

#Data Splitting

In [14]:

```
# Splitting the dependent and independent features
X = data.drop(["Outcome"], axis = 1)
Y = data["Outcome"]

# Splitting the dataset into the training and testing dataset
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.40, random_state=42)

# Printing the size of the training and testing dataset
print("The size of the training dataset: ", X_train.size)
print("The size of the testing dataset: ", X_test.size)
```

The size of the training dataset: 3680

The size of the testing dataset: 2464

In [17]:

```

# Python program to create a function to validate models

def cv_model(models):
    """
    We will create a list of machine learning models and print graphs of cross-validation
    """

    # Cross validating the model using the Kfold stratified cross-validation method
    k_fold = StratifiedKFold(n_splits = 15)

    r = []
    for m in models :
        r.append(cross_val_score(estimator = m, X = X_train, y = Y_train, scoring = "acc

    cross_val_means = []
    cross_val_std = []
    for result in r:
        cross_val_means.append(result.mean())
        cross_val_std.append(result.std())

    df_result = pd.DataFrame({
        "CrossValMean": cross_val_means,
        "CrossValStd": cross_val_std,
        "Model List":[
            "DecisionTreeClassifier",
            "LogisticRegression",
            "SVC",
            "AdaBoostClassifier",
            "GradientBoostingClassifier",
            "RandomForestClassifier",
            "KNeighborsClassifier"
        ]
    })

    # Generating the graph of cross-validation scores
    bar_plot = sns.barplot(x = cross_val_means, y = df_result["Model List"].values, data=
    bar_plot.set_xlabel("Mean of Cross Validation Accuracy Scores")
    bar_plot.set_title("Cross Validation Scores of Models")
    return df_result

```

In [19]:

```
# Importing the required libraries
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV

# Defining a function to analyse the grid results
def analyze_grid(grid):
    """
    Analyzing the results of GridCV method and making predictions for the test data
    Presenting the classification report at the end
    """

    # Printing the best parameter and accuracy score
    print("Tuned hyperparameters: ", grid.best_params_)
    print("Accuracy Score:", grid.best_score_)

    mean_values = grid.cv_results_["mean_test_score"]
    std_values = grid.cv_results_["std_test_score"]
    for m, s, p in zip(mean_values, std_values, grid.cv_results_["params"]):
        print(f"Mean: {m}, Std: {s} * 2, Params: {p}")
        print("The classification Report:")
    Y_true, Y_pred = Y_test, grid.predict(X_test)
    print(classification_report(Y_true, Y_pred))
    print()
```

In [21]:

```

# Defining the Logistic Regression model and its parameters
model = LogisticRegression(solver = 'liblinear')
solver_list = ['liblinear']
penalty_type = ['l2']
c_values = [200, 100, 10, 1.0, 0.01]

# Defining the grid search
grid_lr = dict(solver = solver_list, penalty = penalty_type, C = c_values)
cross_val = StratifiedKFold(n_splits = 100, random_state = 10, shuffle = True)
grid_search_cv = GridSearchCV(estimator = model, param_grid = grid_lr, cv = cross_val, s
lr_result = grid_search_cv.fit(X_train, Y_train)

# Result of Hyper Parameters of Logistic Regression
analyze_grid(lr_result)

```

Tuned hyperparameters: {'C': 200, 'penalty': 'l2', 'solver': 'liblinear'}

Accuracy Score: 0.7715000000000001

Mean: 0.7715000000000001, Std: 0.16556796187668676 * 2, Params: {'C': 200, 'penalty': 'l2', 'solver': 'liblinear'}

The classification Report:

Mean: 0.7715000000000001, Std: 0.16556796187668676 * 2, Params: {'C': 100, 'penalty': 'l2', 'solver': 'liblinear'}

The classification Report:

Mean: 0.7675, Std: 0.16961353129983467 * 2, Params: {'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}

The classification Report:

Mean: 0.7675, Std: 0.17224619008848932 * 2, Params: {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}

The classification Report:

Mean: 0.711, Std: 0.1888888562091475 * 2, Params: {'C': 0.01, 'penalty': 'l2', 'solver': 'liblinear'}

The classification Report:

	precision	recall	f1-score	support
0	0.78	0.88	0.83	201
1	0.70	0.53	0.61	107
accuracy			0.76	308
macro avg	0.74	0.71	0.72	308
weighted avg	0.75	0.76	0.75	308

In [22]:

```
Y_pred = lr_result.predict(X_test)
print(classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.88	0.83	201
1	0.70	0.53	0.61	107
accuracy			0.76	308
macro avg	0.74	0.71	0.72	308
weighted avg	0.75	0.76	0.75	308

In [23]:

```
X_test['predictions'] = Y_pred
print(X_test)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
568	4	154	72	29.000000	126.000000	31.3
620	2	112	86	42.000000	160.000000	38.4
456	1	135	54	20.536458	79.799479	26.7
197	3	107	62	13.000000	48.000000	22.9
714	3	102	74	20.536458	79.799479	29.5
..
70	2	100	66	20.000000	90.000000	32.9
679	2	101	58	17.000000	265.000000	24.2
375	12	140	82	43.000000	325.000000	39.2
700	2	122	76	27.000000	200.000000	35.9
505	10	75	82	20.536458	79.799479	33.3
	DiabetesPedigreeFunction	Age	predictions			
568	0.338	37	1			
620	0.246	28	0			
456	0.687	62	0			
197	0.678	23	0			
714	0.121	32	0			
..			
70	0.867	28	0			
679	0.614	23	0			
375	0.528	58	1			
700	0.483	26	0			
505	0.263	38	0			

[308 rows x 9 columns]